

МОСКОВСКИЙ ГОСУДАРСТВЕННЫЙ ТЕХНИЧЕСКИЙ УНИВЕРСИТЕТ  
им. Н.Э. Баумана

Кафедра «Систем обработки информации и управления»

ОТЧЕТ

**Лабораторная работа №2**  
по курсу «Методы машинного обучения»

Тема: «Изучение библиотек обработки данных»

ИСПОЛНИТЕЛЬ:

Ромичева Е.В.

группа ИУ5-22М

\_\_\_\_\_

подпись

"\_\_" \_\_\_\_\_ 2019 г.

ПРЕПОДАВАТЕЛЬ:

Гапанюк Ю.Е.

\_\_\_\_\_

подпись

"\_\_" \_\_\_\_\_ 2019 г.

Москва - 2019

---

## Цель работы

**Цель лабораторной работы:** изучение различных методов визуализация данных.

## Задание

### Часть 1.

Выполните первое демонстрационное задание "demo assignment" под названием "Exploratory data analysis with Pandas" со страницы курса <https://mlcourse.ai/assignments>

Условие задания -

[https://nbviewer.jupyter.org/github/Yorko/mlcourse\\_open/blob/master/jupyter\\_english/assignments\\_demo/assignment01\\_pandas\\_uci\\_adult.ipynb?flush\\_cache=true](https://nbviewer.jupyter.org/github/Yorko/mlcourse_open/blob/master/jupyter_english/assignments_demo/assignment01_pandas_uci_adult.ipynb?flush_cache=true)

Набор данных можно скачать здесь - <https://archive.ics.uci.edu/ml/datasets/Adult>

Пример решения задания - <https://www.kaggle.com/kashnitsky/a1-demo-pandas-and-uci-adult-dataset-solution>

### Часть 2.

Выполните следующие запросы с использованием двух различных библиотек - [Pandas](#) и [PandaSQL](#):

- один произвольный запрос на соединение двух наборов данных
- один произвольный запрос на группировку набора данных с использованием функций агрегирования

Сравните время выполнения каждого запроса в Pandas и PandaSQL.

#### Часть 1.

### 1. How many men and women (*sex* feature) are represented in this dataset?

```
In [88]: # 1
         data['sex'].value_counts()

Out[88]: Male      21790
         Female    10771
         Name: sex, dtype: int64
```

### 2. What is the average age (*age* feature) of women?

```
In [4]: # 2
         # data.loc[data['sex'] == 'Female', 'age'].mean()
         # data.groupby(['sex'])['age'].mean()
         data['age'][data['sex'] == 'Female'].mean()

Out[4]: 36.85823043357163
```

### 3. What is the percentage of German citizens (*native-country* feature)?

```
In [5]: # 3
         # float((data['native-country'] == 'Germany').sum()) / data.shape[0]
         # data['native-country'].eq('Germany').value_counts(normalize = True)
         100.*data['native-country'].eq('Germany').sum()/data.shape[0]

Out[5]: 0.42074874850281013
```

**4-5. What are the mean and standard deviation of age for those who earn more than 50K per year (*salary* feature) and those who earn less than 50K per year?**

```
In [6]: # 4-5
a = data.groupby(['salary']).agg({'age': 'mean'}).reset_index()
b = data.groupby(['salary']).agg({'age': 'std'}).reset_index()
print("The average age of the rich: {0} +- {1} years, poor - {2} +- {3} years.".format(
    round(a['age'][0]), round(b['age'][0],1),
    round(a['age'][1]), round(b['age'][1],1)))
```

The average age of the rich: 37.0 +- 14.0 years, poor - 44.0 +- 10.5 years.

**6. Is it true that people who earn more than 50K have at least high school education? (*education* – *Bachelors*, *Prof-school*, *Assoc-acdm*, *Assoc-voc*, *Masters* or *Doctorate* feature)**

```
In [7]: deg = {'Bachelors', 'Prof-school', 'Assoc-acdm', 'Assoc-voc', 'Masters', 'Doctorate' }
sal_big = data.loc[data['salary'] == '>50K']
good_edu = sal_big.loc[
    lambda x : x['education'].isin(deg)
]
good_edu['age'].size/sal_big['age'].size
# посчитали отношение образованных людей с хорошей зп и людей с хорошей зп
```

Out[7]: 0.5783701058538452

**7. Display age statistics for each race (*race* feature) and each gender (*sex* feature). Use *groupby()* and *describe()*. Find the maximum age of men of *Amer-Indian-Eskimo* race.**

```
In [8]: # 7
# for (race, sex), sub_df in data.groupby(['race', 'sex']):
#     print("Race: {0}, sex: {1}".format(race, sex))
#     print(sub_df['age'].describe())
sub_df = data.groupby(['race', 'sex'])
sub_df['age'].describe()
```

Out[8]:

		count	mean	std	min	25%	50%	75%	max
race sex									
Amer-Indian-Eskimo	Female	119.0	37.117647	13.114991	17.0	27.0	36.0	46.00	80.0
	Male	192.0	37.208333	12.049563	17.0	28.0	35.0	45.00	82.0
Asian-Pac-Islander	Female	346.0	35.089595	12.300845	17.0	25.0	33.0	43.75	75.0
	Male	693.0	39.073593	12.883944	18.0	29.0	37.0	46.00	90.0
Black	Female	1555.0	37.854019	12.637197	17.0	28.0	37.0	46.00	90.0
	Male	1569.0	37.682600	12.882612	17.0	27.0	36.0	46.00	90.0
Other	Female	109.0	31.678899	11.631599	17.0	23.0	29.0	39.00	74.0
	Male	162.0	34.654321	11.355531	17.0	26.0	32.0	42.00	77.0
White	Female	8642.0	36.811618	14.329093	17.0	25.0	35.0	46.00	90.0
	Male	19174.0	39.652498	13.436029	17.0	29.0	38.0	49.00	90.0

**8. Among whom is the proportion of those who earn a lot (>50K) greater: married or single men (*marital-status* feature)? Consider as married those who have a *marital-status* starting**

with *Married* (Married-civ-spouse, Married-spouse-absent or Married-AF-spouse), the rest are considered bachelors.

```
In [11]: # 8
mar = {'Married-civ-spouse', 'Married-spouse-absent ', 'Married-AF-spouse'}
men = data.loc[data['sex'] == 'Male']
sal_big = men.loc[data['salary'] == '>50K']
married_salary = sal_big.loc[
    lambda x : x['marital-status'].isin(mar)
]
print ('Отношение богатых замужних мужчин к богатым незамужним = ', married_salary['age'].size/sal_big['age'].size)
sal_big['marital-status'].value_counts()

Отношение богатых замужних мужчин к богатым незамужним =  0.89192434704293

Out[11]: Married-civ-spouse      5938
Never-married      325
Divorced      284
Separated      49
Widowed      39
Married-spouse-absent      23
Married-AF-spouse      4
Name: marital-status, dtype: int64
```

9. What is the maximum number of hours a person works per week (*hours-per-week* feature)? How many people work such a number of hours, and what is the percentage of those who earn a lot (>50K) among them?

```
In [14]: # 9
max_load = data['hours-per-week'].max()
print("Max time - {0} hours./week.".format(max_load))
num_workaholics = data[data['hours-per-week'] == max_load].shape[0]
print("Total number of such hard workers {0}".format(num_workaholics))
rich_workaholics = float(data[(data['hours-per-week'] == max_load)
    & (data['salary'] == '>50K')].shape[0]) / num_workaholics
print("Percentage of rich among them {0}%".format(int(100 * rich_workaholics)))

Max time - 99 hours./week.
Total number of such hard workers 85
Percentage of rich among them 29%
```

10. Count the average time of work (*hours-per-week*) for those who earn a little and a lot (*salary*) for each country (*native-country*). What will these be for Japan?

```
In [13]: # 10
zz = data.groupby(['native-country', 'salary']).agg({'hours-per-week': 'mean'})
zz
```

Out[13]:

		hours-per-week	
native-country salary			
?	<=50K	40.164760	
	>50K	45.547945	
Cambodia	<=50K	41.416667	
	>50K	40.000000	
Canada	<=50K	37.914634	
	>50K	45.641026	
China	<=50K	37.381818	
	>50K	38.900000	
Columbia	<=50K	38.684211	
	>50K	50.000000	
Cuba	<=50K	37.985714	
	>50K	42.440000	
Dominican-Republic	<=50K	42.338235	
	>50K	47.000000	
Ecuador	<=50K	38.041667	
	>50K	48.750000	
El-Salvador	<=50K	36.030928	
	>50K	45.000000	
England	<=50K	40.483333	
	>50K	44.533333	
France	<=50K	41.058824	
	>50K	50.750000	

## Часть 2.

### Соединение таблиц с помощью Pandas

```
In [9]: def example1_pandas(project_submissions, enrollments):
merg_res = pd.merge(project_submissions, enrollments, on='account_key')
return merg_res

example1_pandas(project_submissions, enrollments)
```

Out[9]:

	creation_date	completion_date	assigned_rating	account_key	lesson_key	processing_state	status	join_date	cancel_date	days_to_cancel	is_udacity
0	2015-01-14	2015-01-16	UNGRADED	256	3176718735	EVALUATED	canceled	2014-12-03	2015-04-01	119.0	False
1	2015-01-14	2015-01-16	UNGRADED	256	3176718735	EVALUATED	canceled	2015-04-01	2015-06-10	70.0	False
2	2015-01-10	2015-01-13	INCOMPLETE	256	3176718735	EVALUATED	canceled	2014-12-03	2015-04-01	119.0	False
3	2015-01-10	2015-01-13	INCOMPLETE	256	3176718735	EVALUATED	canceled	2015-04-01	2015-06-10	70.0	False
4	2015-01-20	2015-01-20	PASSED	256	3176718735	EVALUATED	canceled	2014-12-03	2015-04-01	119.0	False
5	2015-01-20	2015-01-20	PASSED	256	3176718735	EVALUATED	canceled	2015-04-01	2015-06-10	70.0	False
6	2015-03-10	2015-03-13	PASSED	434	3176718735	EVALUATED	canceled	2015-01-12	2015-06-03	142.0	False

### Соединение таблиц с помощью PandaSQL

```
In [12]: pysql = lambda q: ps.sqldf(q, globals())
def example1_pandasql(project_submissions, enrollments):
#     query = """SELECT *
#     FROM project_submissions p JOIN enrollments e ON p.account_key = e.account_key
#     GROUP BY p.account_key;"""
    query = "select * from project_submissions, enrollments where project_submissions.account_key = enrollments.account_key;"
    join_res = pysql(query)
    return join_res

example1_pandasql(project_submissions, enrollments)
```

```
Out[12]:
```

	creation_date	completion_date	assigned_rating	account_key	lesson_key	processing_state	account_key	status	join_date	cancel_date	days_to_c
0	2015-01-14	2015-01-16	UNGRADED	256	3176718735	EVALUATED	256	canceled	2014-12-03	2015-04-01	
1	2015-01-14	2015-01-16	UNGRADED	256	3176718735	EVALUATED	256	canceled	2015-04-01	2015-06-10	
2	2015-01-10	2015-01-13	INCOMPLETE	256	3176718735	EVALUATED	256	canceled	2014-12-03	2015-04-01	
3	2015-01-10	2015-01-13	INCOMPLETE	256	3176718735	EVALUATED	256	canceled	2015-04-01	2015-06-10	
4	2015-01-20	2015-01-20	PASSED	256	3176718735	EVALUATED	256	canceled	2014-12-03	2015-04-01	
5	2015-01-20	2015-01-20	PASSED	256	3176718735	EVALUATED	256	canceled	2015-04-01	2015-06-10	
6	2015-03-10	2015-03-13	PASSED	434	3176718735	EVALUATED	434	canceled	2015-01-12	2015-06-03	

## Сравнение времени выполнения запросов

```
In [10]: class Profiler(object):
def __enter__(self):
    self._startTime = time.time()

def __exit__(self, type, value, traceback):
    print ("Elapsed time: {:.3f} sec".format(time.time() - self._startTime))

with Profiler() as p:
    example1_pandas(project_submissions, enrollments)

Elapsed time: 0.009 sec
```

```
In [11]: with Profiler() as p:
example1_pandasql(project_submissions, enrollments)

Elapsed time: 0.105 sec
```

Видно, что соединение таблиц с помощью Pandas выполняется в 12 раз быстрее, чем соединение с помощью PandaSQL.

## Агрегирование таблиц с помощью PandaSQL

```
In [12]: daily_engagements['weekday'] = list(map(lambda x: datetime.strptime(x, '%Y-%m-%d').strftime('%A'),
daily_engagements.utc_date))

# pandasql code
def example2_pandasql(daily_engagements):
    aggr_query = '''
        SELECT
            avg(total_minutes_visited) as total_minutes_visited,
            weekday
        FROM daily_engagements
        GROUP BY weekday
    '''
    return ps.sqldf(aggr_query, locals()).set_index('weekday')

# pandas code
def example2_pandas(daily_engagements):
    return pd.DataFrame(daily_engagements.groupby('weekday').total_minutes_visited.mean())

weekday_engagement = example2_pandasql(daily_engagements)
weekday_engagement
```

Out[12]:

total_minutes_visited	
weekday	
Friday	23.156233
Monday	26.418982
Saturday	21.725677
Sunday	23.539406
Thursday	24.685176
Tuesday	26.857676
Wednesday	25.362789

Агрегирование таблиц с помощью Pandas

```
In [13]: example2_pandas(daily_engagements)
```

Out[13]:

total_minutes_visited	
weekday	
Friday	23.156233
Monday	26.418982
Saturday	21.725677
Sunday	23.539406
Thursday	24.685176
Tuesday	26.857676
Wednesday	25.362789

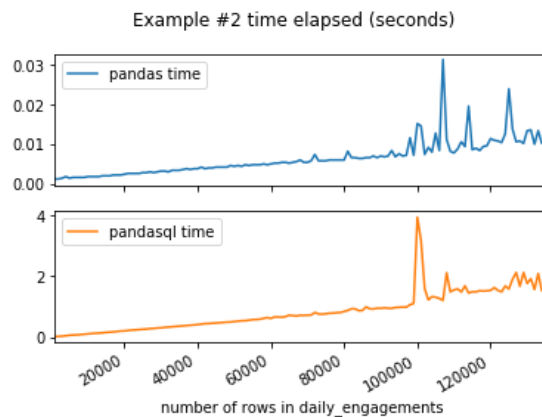
Сравнение времени выполнения запросов

```
In [5]: ex2_times = []
        for count in range(1000, 137000, 1000):
            pandasql_time = count_mean_time(example2_pandasql, [daily_engagements[:count]])
            pandas_time = count_mean_time(example2_pandas, [daily_engagements[:count]])
            ex2_times.append({'count': count, 'pandasql_time': pandasql_time, 'pandas_time': pandas_time})
```

```
In [6]: ex2_times_df = pd.DataFrame(ex2_times)
```

```
In [7]: ex2_times_df.columns = ['number of rows in daily_engagements', 'pandas time', 'pandasql time']
        ex2_times_df = ex2_times_df.set_index('number of rows in daily_engagements')
```

```
In [8]: ax = ex2_times_df.plot(title = 'Example #2 time elapsed (seconds)', subplots = True)
```



Видно, что агрегирование таблиц с помощью Pandas выполняется быстрее, чем агрегирование с помощью PandaSQL.