

МОСКОВСКИЙ ГОСУДАРСТВЕННЫЙ ТЕХНИЧЕСКИЙ УНИВЕРСИТЕТ
им. Н.Э. Баумана

Кафедра «Систем обработки информации и управления»

ОТЧЕТ

Лабораторная работа №4
по курсу «Методы машинного обучения»

Тема: «Подготовка обучающей и тестовой выборки, кросс-валидация и подбор гиперпараметров на примере метода ближайших соседей.»

ИСПОЛНИТЕЛЬ:

Ромичева Е.В.

группа ИУ5-22М

подпись

"__" _____ 2019 г.

ПРЕПОДАВАТЕЛЬ:

Гапанюк Ю.Е.

подпись

"__" _____ 2019 г.

Москва - 2019

Цель работы

Цель лабораторной работы: изучение сложных способов подготовки выборки и подбора гиперпараметров на примере метода ближайших соседей.

Задание

1. Выберите набор данных (датасет) для решения задачи классификации или регрессии.
2. В случае необходимости проведите удаление или заполнение пропусков и кодирование категориальных признаков.
3. С использованием метода `train_test_split` разделите выборку на обучающую и тестовую.
4. Обучите модель ближайших соседей для произвольно заданного гиперпараметра K . Оцените качество модели с помощью трех подходящих для задачи метрик.
5. Постройте модель и оцените качество модели с использованием кросс-валидации. Проведите эксперименты с тремя различными стратегиями кросс-валидации.
6. Произведите подбор гиперпараметра K с использованием `GridSearchCV` и кросс-валидации.
7. Повторите пункт 4 для найденного оптимального значения гиперпараметра K . Сравните качество полученной модели с качеством модели, полученной в пункте 4.
8. Постройте кривые обучения и валидации.

lab4

March 24, 2019

```
In [1]: # https://www.kaggle.com/uciml/pima-indians-diabetes-database
# https://www.kaggle.com/amolbhivarkar/knn-for-classification-using-scikit-learn

In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import learning_curve, validation_curve
from sklearn.model_selection import KFold, RepeatedKFold, LeaveOneOut, LeavePOut, ShuffleSplit
from sklearn.model_selection import cross_val_score, cross_validate
from sklearn.metrics import roc_curve, confusion_matrix, roc_auc_score, accuracy_score,
plt.style.use('ggplot')
```

```
In [2]: #Load the dataset
df = pd.read_csv('diabetes.csv')

#Print the first 5 rows of the dataframe.
df.head()
```

Out[2]:	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
In [6]: df.shape

Out[6]: (768, 9)

In [7]: for col in df.columns:
#     -
```

```

        temp_null_count = df[df[col].isnull()].shape[0]
        print('{} - {}'.format(col, temp_null_count))

Pregnancies - 0
Glucose - 0
BloodPressure - 0
SkinThickness - 0
Insulin - 0
BMI - 0
DiabetesPedigreeFunction - 0
Age - 0
Outcome - 0

In [3]: X = df.drop('Outcome',axis=1).values
        y = df['Outcome'].values

In [9]: from sklearn.neighbors import KNeighborsClassifier
        from sklearn.model_selection import train_test_split

In [10]: X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.4,random_state=42, s

        #Setup arrays to store training and test accuracies
        neighbors = np.arange(1,9)
        train_accuracy =np.empty(len(neighbors))
        test_accuracy = np.empty(len(neighbors))

        for i,k in enumerate(neighbors):
            #Setup a knn classifier with k neighbors
            knn = KNeighborsClassifier(n_neighbors=k)


            #Fit the model
            knn.fit(X_train, y_train)

            #Compute accuracy on the training set
            train_accuracy[i] = knn.score(X_train, y_train)

            #Compute accuracy on the test set
            test_accuracy[i] = knn.score(X_test, y_test)

In [11]: #Generate plot
        plt.title('k-NN Varying number of neighbors')
        plt.plot(neighbors, test_accuracy, label='Testing Accuracy')
        plt.plot(neighbors, train_accuracy, label='Training accuracy')
        plt.legend()
        plt.xlabel('Number of neighbors')
        plt.ylabel('Accuracy')
        plt.show()

```



```
In [12]: #Setup a knn classifier with k neighbors
knn = KNeighborsClassifier(n_neighbors=7)
```

```
In [13]: #Fit the model
knn.fit(X_train,y_train)
```

```
Out[13]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
metric_params=None, n_jobs=None, n_neighbors=7, p=2,
weights='uniform')
```

```
In [14]: #Get accuracy. Note: In case of classification algorithms score method represents acc
knn.score(X_test,y_test)
```

```
Out[14]: 0.7305194805194806
```

```
In [16]: #import classification_report
from sklearn.metrics import classification_report
y_pred = knn.predict(X_test)
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.78	0.82	0.80	201
1	0.62	0.56	0.59	107

micro avg	0.73	0.73	0.73	308
macro avg	0.70	0.69	0.70	308
weighted avg	0.73	0.73	0.73	308

0.0.1 Accuracy

```
In [29]: # 7
         cl1_1 = KNeighborsClassifier(n_neighbors=7)
         cl1_1.fit(X_train, y_train)
         target1_1 = cl1_1.predict(X_test)
         accuracy_score(y_test, target1_1)
```

```
Out[29]: 0.7305194805194806
```


0.0.2 Confusion Matrix

```
In [21]: y_pred = knn.predict(X_test)
         confusion_matrix(y_test,y_pred)
         pd.crosstab(y_test, y_pred, rownames=['True'], colnames=['Predicted'], margins=True)
```

```
Out[21]: Predicted    0    1  All
         True
         0          165   36  201
         1           47   60  107
         All         212   96  308
```

0.0.3 ROC (Reciever Operating Charecteristic) curve

```
In [24]: y_pred_proba = knn.predict_proba(X_test)[: ,1]
         fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
         plt.plot([0,1],[0,1], 'k--')
         plt.plot(fpr,tpr, label='Knn')
         plt.xlabel('fpr')
         plt.ylabel('tpr')
         plt.title('Knn(n_neighbors=7) ROC curve')
         plt.show()
```



```
In [27]: roc_auc_score(y_test,y_pred_proba)
```

```
Out[27]: 0.7345050448691124
```

0.1 Cross Validation

```
In [31]: param_grid = {'n_neighbors':np.arange(1,50)}
knn = KNeighborsClassifier()
knn_cv= GridSearchCV(knn,param_grid,cv=5)
knn_cv.fit(X,y)
```

```
Out[31]: GridSearchCV(cv=5, error_score='raise-deprecating',
    estimator=KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
    metric_params=None, n_jobs=None, n_neighbors=5, p=2,
    weights='uniform'),
    fit_params=None, iid='warn', n_jobs=None,
    param_grid={'n_neighbors': array([ 1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49])},
    pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
    scoring=None, verbose=0)
```

```
In [32]: knn_cv.best_score_
```

```
Out [32]: 0.7578125
```

```
In [33]: knn_cv.best_params_
```

```
Out [33]: {'n_neighbors': 14}
```

0.1.1 K-fold

```
In [38]: scores = cross_val_score(KNeighborsClassifier(n_neighbors=2),  
                                   X, y,  
                                   cv=KFold(n_splits=3))  
  
scores
```

```
Out [38]: array([0.671875 , 0.72265625, 0.73046875])
```

0.1.2 Leave One Out (LOO)

```
In [1]: # # KFold(n_splits=n)  
        # loo = LeaveOneOut()  
        # loo.get_n_splits(X)  
  
        # for train_index, test_index in loo.split(X):  
        #     print("TRAIN:", train_index, "TEST:", test_index)  
        #     X_train, X_test = X[train_index], X[test_index]  
        #     y_train, y_test = y[train_index], y[test_index]  
        #     print(X_train, X_test, y_train, y_test)
```

0.1.3 Repeated K-Fold

```
In [39]: scores2 = cross_val_score(KNeighborsClassifier(n_neighbors=2),  
                                   X, y,  
                                   cv=RepeatedKFold(n_splits=3, n_repeats=2))  
  
scores2
```

```
Out [39]: array([0.70703125, 0.70703125, 0.6953125 , 0.69140625, 0.6875  
                0.75      ])
```

0.2 K

```
In [41]: knn = KNeighborsClassifier(n_neighbors=14)  
        knn.fit(X_train,y_train)  
        knn.score(X_test,y_test)
```

```
Out [41]: 0.7305194805194806
```

```
In [42]: # accuracy  
        cl1_1 = KNeighborsClassifier(n_neighbors=14)  
        cl1_1.fit(X_train, y_train)  
        target1_1 = cl1_1.predict(X_test)  
        accuracy_score(y_test, target1_1)
```

```
Out [42]: 0.7305194805194806
```


0.2.1

```
In [46]: def plot_learning_curve(estimator, title, X, y, ylim=None, cv=None,
                                n_jobs=None, train_sizes=np.linspace(.1, 1.0, 5)):

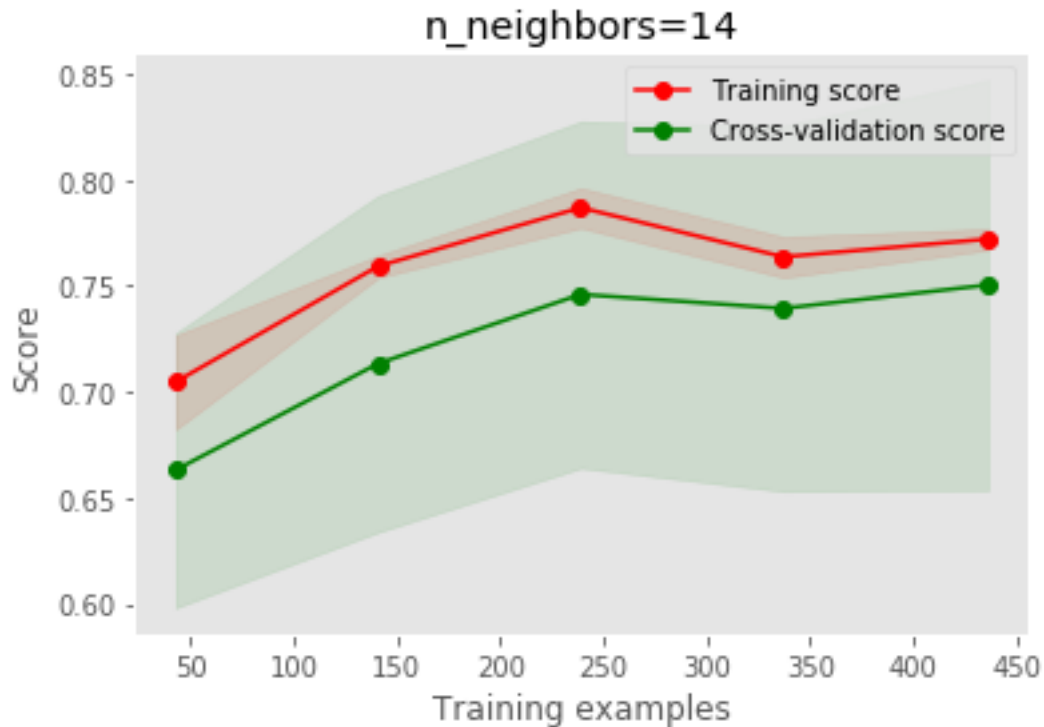
    plt.figure()
    plt.title(title)
    if ylim is not None:
        plt.ylim(*ylim)
    plt.xlabel("Training examples")
    plt.ylabel("Score")
    train_sizes, train_scores, test_scores = learning_curve(
        estimator, X, y, cv=cv, n_jobs=n_jobs, train_sizes=train_sizes)
    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)
    test_scores_mean = np.mean(test_scores, axis=1)
    test_scores_std = np.std(test_scores, axis=1)
    plt.grid()

    plt.fill_between(train_sizes, train_scores_mean - train_scores_std,
                     train_scores_mean + train_scores_std, alpha=0.1,
                     color="r")
    plt.fill_between(train_sizes, test_scores_mean - test_scores_std,
                     test_scores_mean + test_scores_std, alpha=0.1, color="g")
    plt.plot(train_sizes, train_scores_mean, 'o-', color="r",
             label="Training score")
    plt.plot(train_sizes, test_scores_mean, 'o-', color="g",
             label="Cross-validation score")

    plt.legend(loc="best")
    return plt

In [47]: plot_learning_curve(KNeighborsClassifier(n_neighbors=14), 'n_neighbors=14',
                             X_train, y_train, cv=20)
```

```
Out[47]: <module 'matplotlib.pyplot' from 'c:\\users\\helen\\appdata\\local\\programs\\python\\
```



0.2.2

```
In [48]: def plot_validation_curve(estimator, title, X, y,
                                   param_name, param_range, cv,
                                   scoring="accuracy"):

    train_scores, test_scores = validation_curve(
        estimator, X, y, param_name=param_name, param_range=param_range,
        cv=cv, scoring=scoring, n_jobs=1)
    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)
    test_scores_mean = np.mean(test_scores, axis=1)
    test_scores_std = np.std(test_scores, axis=1)

    plt.title(title)
    plt.xlabel(param_name)
    plt.ylabel("Score")
    plt.ylim(0.0, 1.1)
    lw = 2
    plt.plot(param_range, train_scores_mean, label="Training score",
             color="darkorange", lw=lw)
    plt.fill_between(param_range, train_scores_mean - train_scores_std,
                    train_scores_mean + train_scores_std, alpha=0.2,
```

```

        color="darkorange", lw=lw)
plt.plot(param_range, test_scores_mean, label="Cross-validation score",
        color="navy", lw=lw)
plt.fill_between(param_range, test_scores_mean - test_scores_std,
        test_scores_mean + test_scores_std, alpha=0.2,
        color="navy", lw=lw)
plt.legend(loc="best")
return plt

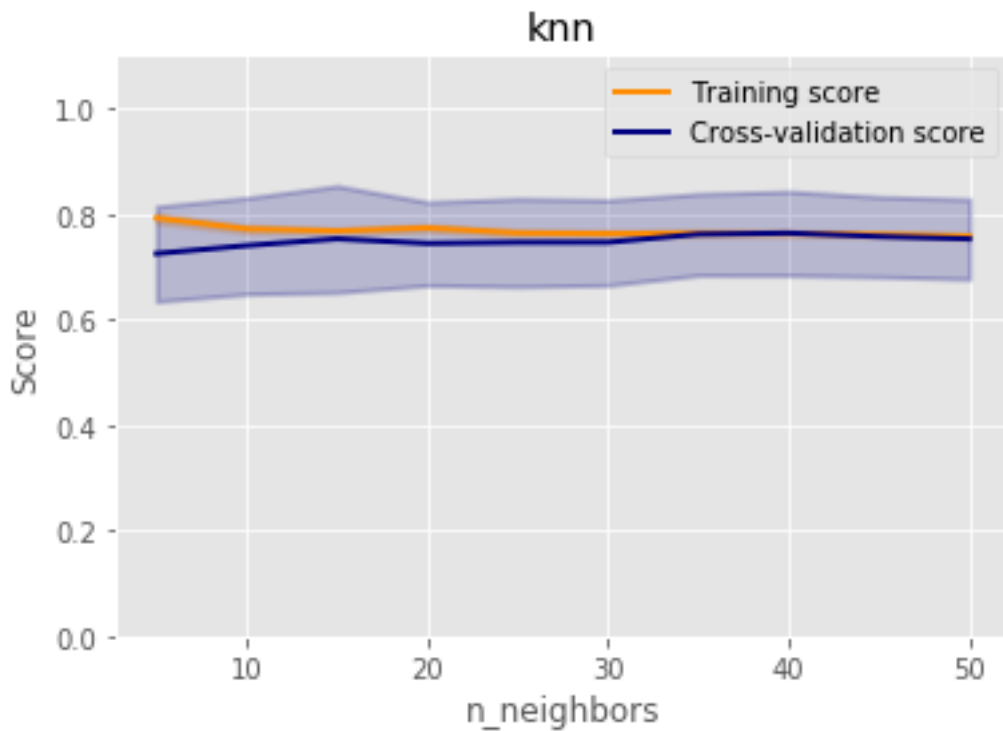
```

```

In [50]: n_range = np.array(range(5,55,5))
        plot_validation_curve(KNeighborsClassifier(), 'knn',
                               X_train, y_train,
                               param_name='n_neighbors', param_range=n_range,
                               cv=20, scoring="accuracy")

```

Out[50]: <module 'matplotlib.pyplot' from 'c:\\users\\helen\\appdata\\local\\programs\\python\\



In []: